

Program Flow (based on user-driven events):

- Event: Starting a new chat
 - Start listening as leader, initializing
- Event: Node N joining a chat
 - From N 's perspective, given a known member M of the group:
 - Loop:
 - Send M an ADD_ME message.
 - Wait until M replies:
 - If the reply is a PARTICIPANT_UPDATE message, the join was successful.
 - If the reply is a JOIN_FAILURE message, wait a short period
 - If the reply is a LEADER_ID message:
 - Send the leader an ADD_ME message
 - If the reply is a PARTICIPANT_UPDATE message, the join was successful.
 - If the reply is a JOIN_FAILURE or LEADER_ID message, wait a short period
 - If the reply is a JOIN_NICKNAME_FAILURE, exit and prompt the user to choose a new nickname
 - After 10 failed attempts, exit program
 - When a member of the group M is contacted by a new node N :
 - If M is leader:
 - If M 's name is in use, send JOIN_NICKNAME_FAILURE message
 - Otherwise, M adds N to list
 - M multicasts new list of nodes to network, including N (PARTICIPANT_UPDATE message)
 - If M is not leader:
 - M sends "Leader is L" message to N
 - If M is in election:
 - M sends N JOIN_FAILURE
 - Wait for an announcement from the leader that N has been added
 - Reply to N with the location of the leader
 - If in leader election or request to leader triggers election, reply to N with JOIN_FAILURE message
- Event: Exchanging messages
 - When a node N wants to send a message m :
 - N sends the leader the message m
 - If a leader election begins, place m in a send queue

- When the leader receives a message m :
 - It assigns the message its current (logical) clock value t and increments the clock
 - For each member of the group:
 - Send the member the tuple $\langle m, t \rangle$
- When a non-leader A receives a message from the leader L
 - Print it
- When a node receives a START_ELECTION message from N :
 - Compare own nickname to that of N
 - If own name is higher:
 - Reply to N with an ELECTION_STOP message.
 - Follow procedure for starting a new election.
- Event: A node N leaves the chat
 - From N 's perspective:
 - Just exit.
 - If N is not a leader:
 - From the leader's Perspective:
 - Tries to send a message to N
 - Message times out
 - Leaders announces N has left (with a PARTICIPANT_UPDATE message)
 - From everyone else's perspective:
 - Gets PARTICIPANT_UPDATE message and updates local state
 - If N is a leader:
 - A node M will try to send N a message m
 - Message times out
 - M broadcasts a START_ELECTION message to the group
 - If M receives an ELECTION_STOP message from a peer, wait until an ELECTION_VICTORY message is received and re-send the message m to the new leader. ($M \rightarrow \text{leader}$)
 - If M receives an ELECTION_STOP message but times out waiting for an ELECTION_VICTORY, resend the START_ELECTION message.
 - If nobody replies to the START_ELECTION message after a timeout, broadcast an ELECTION_VICTORY message.
- Heartbeats?

Node State (For a node M):

- Nickname ($M \rightarrow \text{nickname}$)
- IP address ($M \rightarrow \text{ip_address}$)
- Port # ($M \rightarrow \text{port_num}$)
- Logical clock (sequence number for each message) ($M \rightarrow \text{seq_num}$)
- List of other nodes ($M \rightarrow \text{participants}$)

- Is Leader (*M->is_leader*)
- Reference to leader (*M->leader*)
- Message queue (*M->msg_queue*)

Message Types:

- PARTICIPANT_UPDATE:
 - Contains a list of rows of the following form:
 - <nickname, IP address, Port #, is_leader>

Transport Layer: Reliable Message Protocol:

- A datagram-oriented protocol that guarantees that messages will be delivered exactly once.
- `int RMP_get_address_for(const char *address, int port, rmp_address *address_dst);`
 - Parses the given IP address and port number into the `rmp_address` in `*address_dst`.
 - Returns 0 on success and -1 on error.
- `int RMP_sendTo(rmp_address address, const void *buffer, int numBytes);`
 - Attempts to send `numBytes` bytes of the specified `buffer` to the given `address`.
 - Will send up to 65,502 bytes of data at a time, all remaining bytes will be truncated.
 - Will block and retry sending until the recipient acknowledges the message OR at least the timeout value of 100ms is reached.
 - Returns the number of bytes sent or -1 on error.
- `int RMP_setup_receiving_socket(rmp_address address);`
 - Sets up an RMP socket that can be listened on for incoming connections on the given `address`.
 - Returns a socket file descriptor or -1 on error.
- `int RMP_listen(int socket_fd, void *buffer, size_t len, rmp_address *src_addr);`
 - Blocks for a message to be received on `socket_fd` and places up to `len` bytes of it in the given `buffer`. Also places the sender's address in `*src_addr`.
 - Returns the number of bytes received or -1 on error.
- Implemented over UDP with the following packet format:
 - | | | |
|--------|-----------|------|
| Type | ID number | Data |
| 1 byte | 4 bytes | ... |
 - Type is one of DATA, SYN_ACK, or ACK
- When a node A wishes to send a message to another node B, the following protocol is used:

- When A wants to send B a message, it first constructs a `DATA` packet with the desired message and a random ID number. It sends the message to B and stores a duplicate in a send buffer.
- When B receives a `DATA` packet from A, it stores the ID number in a receive buffer for A and replies to A with an `ACK` packet containing the ID.
- When A receives an `ACK` packet from B, it replies with a `SYN_ACK` packet with the same ID number. If the ID number matches the one in its send buffer, it removes it from the buffer and becomes ready to send again.
- When B receives a `SYN_ACK` packet from A, it removes the corresponding ID from its receive buffer.
- If any item remains in a send or receive buffer for more than the time (`TIMEOUT/NUM_RETRIES`), its corresponding message is retransmitted. After it has been retransmitted `NUM_RETRIES` number of times, it is removed.